



**KOÇ
UNIVERSITY**

Database Management Systems

Structured Query Language (SQL)

M. Emre Gürsoy



Assistant Professor
Department of Computer Engineering

www.memregursoy.com



Introduction

- Pronounced as "S-Q-L" or "SeQueL"
- SQL is considered to be one of the major reasons for the commercial success of the **relational** data model
 - IBM, Microsoft, Oracle, ...
 - Main push from IBM in the early days (1970s)





Some Important Syntax

- **CREATE DATABASE** *databasename*;
- **CREATE TABLE** *table_name* (
 column1 datatype,
 column2 datatype,
 column3 datatype,

);
- **INSERT, UPDATE, DELETE**
- **SELECT** *column1, column2, ...*
 FROM *table_name*
 WHERE *condition*;
- **DROP TABLE** *table_name*

Name	Date of Birth	Gender	Zipcode	Disease
Andre	21/1/76	Male	53715	Heart Disease
Beth	13/4/86	Female	53715	Hepatitis
Carol	28/2/76	Female	53703	Brochitis
Dan	9/3/92	Male	53703	Broken Arm
Ellen	2/6/88	Female	53706	Flu
Erica	28/12/97	Female	53706	Hang Nail



Creating Tables

- **CREATE TABLE EMPLOYEE . . .**
 - We have been working in "pseudo"-SQL
- Columns have **data types**
 - Numeric: `INTEGER`, `FLOAT`, `REAL`, ...
 - String: `CHAR (n)`, `CHARACTER (n)`, `VARCHAR (n)`
 - Bit-string: `BIT (n)`
 - Boolean: `TRUE` **or** `FALSE` **or** `NULL`
 - Date: `YEAR`, `MONTH`, and `DAY` in the form `YYYY-MM-DD`
 - Timestamp: includes date and time
 - ...
- **UNIQUE** clause: Used to represent a CANDIDATE key that was not selected as the PRIMARY key



Company DB

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------

DEPT_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
----------------	------------------

PROJECT

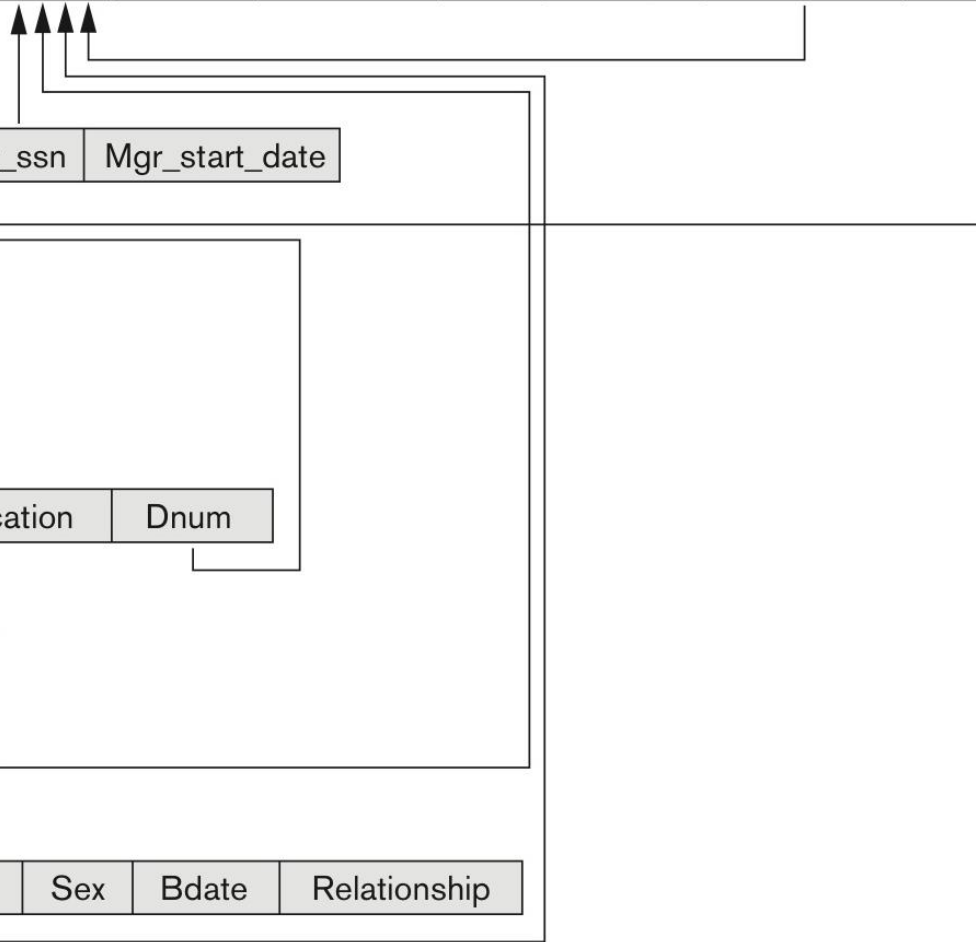
Pname	<u>Pnumber</u>	Plocation	Dnum
-------	----------------	-----------	------

WORKS_ON

<u>Essn</u>	<u>Pno</u>	Hours
-------------	------------	-------

DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
-------------	-----------------------	-----	-------	--------------





Company DB

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

DEPT_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
1	Houston
4	Stafford
5	Bellaire
5	Sugarland
5	Houston



Company DB

WORKS_ON

<u>Essn</u>	<u>Pno</u>	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0
987654321	20	15.0
888665555	20	NULL

PROJECT

<u>Pname</u>	<u>Pnumber</u>	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
333445555	Alice	F	1986-04-05	Daughter
333445555	Theodore	M	1983-10-25	Son
333445555	Joy	F	1958-05-03	Spouse
987654321	Abner	M	1942-02-28	Spouse
123456789	Michael	M	1988-01-04	Son
123456789	Alice	F	1988-12-30	Daughter
123456789	Elizabeth	F	1967-05-05	Spouse



Company DB

CREATE TABLE EMPLOYEE

(Fname	VARCHAR(15)	NOT NULL,
Minit	CHAR,	
Lname	VARCHAR(15)	NOT NULL,
Ssn	CHAR(9)	NOT NULL,
Bdate	DATE,	
Address	VARCHAR(30),	
Sex	CHAR,	
Salary	DECIMAL(10,2),	
Super_ssn	CHAR(9),	
Dno	INT	NOT NULL,

PRIMARY KEY (Ssn),

CREATE TABLE DEPARTMENT

(Dname	VARCHAR(15)	NOT NULL,
Dnumber	INT	NOT NULL,
Mgr_ssn	CHAR(9)	NOT NULL,
Mgr_start_date	DATE,	

PRIMARY KEY (Dnumber),

UNIQUE (Dname),

FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn));

CREATE TABLE DEPT_LOCATIONS

(Dnumber	INT	NOT NULL,
Dlocation	VARCHAR(15)	NOT NULL,

PRIMARY KEY (Dnumber, Dlocation),

FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber));



Company DB

CREATE TABLE PROJECT

(Pname	VARCHAR(15)	NOT NULL,
Pnumber	INT	NOT NULL,
Plocation	VARCHAR(15),	
Dnum	INT	NOT NULL,

PRIMARY KEY (Pnumber),

UNIQUE (Pname),

FOREIGN KEY (Dnum) **REFERENCES** DEPARTMENT(Dnumber));

CREATE TABLE WORKS_ON

(Essn	CHAR(9)	NOT NULL,
Pno	INT	NOT NULL,
Hours	DECIMAL(3,1)	NOT NULL,

PRIMARY KEY (Essn, Pno),

FOREIGN KEY (Essn) **REFERENCES** EMPLOYEE(Ssn),

FOREIGN KEY (Pno) **REFERENCES** PROJECT(Pnumber));

CREATE TABLE DEPENDENT

(Essn	CHAR(9)	NOT NULL,
Dependent_name	VARCHAR(15)	NOT NULL,
Sex	CHAR,	
Bdate	DATE,	
Relationship	VARCHAR(8),	

PRIMARY KEY (Essn, Dependent_name),

FOREIGN KEY (Essn) **REFERENCES** EMPLOYEE(Ssn));



Retrieval Queries

SELECT	[DISTINCT] <i>attribute-list</i>
FROM	<i>relation-list</i>
WHERE	<i>condition</i>

- **attribute-list** is a list of attribute names whose values are to be retrieved by the query.
- **relation-list** is a list of relation names (table names) required to process the query.
- **condition** is a boolean expression that specifies what conditions the tuples need to satisfy to be retrieved.
- **DISTINCT** is an optional keyword indicating duplicates should be eliminated (otherwise no duplicate elimination).



Retrieval Queries

<u>Bdate</u>	<u>Address</u>
1965-01-09	731 Fondren, Houston, TX

<u>Fname</u>	<u>Lname</u>	<u>Address</u>
John	Smith	731 Fondren, Houston, TX
Franklin	Wong	638 Voss, Houston, TX
Ramesh	Narayan	975 Fire Oak, Humble, TX
Joyce	English	5631 Rice, Houston, TX

Query 0. Retrieve the birth date and address of the employee(s) whose name is 'John B. Smith'.

Q0: **SELECT** Bdate, Address
 FROM EMPLOYEE
 WHERE Fname='John' **AND** Minit='B' **AND** Lname='Smith';

Query 1. Retrieve the name and address of all employees who work for the 'Research' department.

Q1: **SELECT** Fname, Lname, Address
 FROM EMPLOYEE, DEPARTMENT
 WHERE Dname='Research' **AND** Dnumber=Dno;



Retrieval Queries

(c)

<u>Pnumber</u>	<u>Dnum</u>	<u>Lname</u>	<u>Address</u>	<u>Bdate</u>
10	4	Wallace	291Berry, Bellaire, TX	1941-06-20
30	4	Wallace	291Berry, Bellaire, TX	1941-06-20

Query 2. For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birth date.

Q2: **SELECT** Pnumber, Dnum, Lname, Address, Bdate
 FROM PROJECT, DEPARTMENT, EMPLOYEE
 WHERE Dnum=Dnumber **AND** Mgr_ssn=Ssn **AND**
 Plocation='Stafford';



Attribute Naming

- If attribute names are unique across the referenced relations, we can omit the relation name in the query
 - We can write **Plocation = "Stafford"** instead of **Project.Plocation = "Stafford"**
- But if there exist two attributes with the same name in different relations, we must write the relation explicitly
 - In order to prevent ambiguity
- Some people prefer to always write the relation name

```
Q1A:  SELECT  Fname, EMPLOYEE.Name, Address
        FROM    EMPLOYEE, DEPARTMENT
        WHERE   DEPARTMENT.Name='Research' AND
                DEPARTMENT.Dnumber=EMPLOYEE.Dnumber;
```



Aliasing and Renaming

- **AS**: Declare alternative names for relations
 - **EMPLOYEE AS E, DEPARTMENT AS D, ...**
 - This strategy is sometimes used just for brevity, but sometimes it becomes necessary
- Keyword **AS** can also be dropped

Query 8. For each employee, retrieve the employee's first and last name and the last name of the employee's supervisor.

```
SELECT      E.Fname, E.Lname, S.Lname
FROM        EMPLOYEE AS E, EMPLOYEE AS S
WHERE       E.Super_ssn=S.Ssn;
```



Missing "Where" Clause

- Indicates no condition on tuple selection
- Effect is a CROSS PRODUCT
 - Recall from relational algebra

Queries 9 and 10. Select all EMPLOYEE Ssns (Q9) and all combinations of EMPLOYEE Ssn and DEPARTMENT Dname (Q10) in the database.

Q9: **SELECT** Ssn
 FROM EMPLOYEE;

Q10: **SELECT** Ssn, Dname
 FROM EMPLOYEE, DEPARTMENT;



Use of Asterisk (*)

- Retrieve **all** attribute values of selected tuples
- Effect is "having no projection" (in relational algebra terms)

Q1C: **SELECT** *

FROM EMPLOYEE

WHERE Dno=5;

Q1D: **SELECT** *

FROM EMPLOYEE, DEPARTMENT

WHERE Dname='Research' **AND** Dno=Dnumber;

Q10A: **SELECT** *

FROM EMPLOYEE, DEPARTMENT;



DISTINCT Keyword

- SQL does not automatically eliminate duplicates
 - This is different from the default behavior of the projection operation in relational algebra
- We use the **DISTINCT** keyword to eliminate duplicates

Query 11. Retrieve the salary of every employee (Q11) and all distinct salary values (Q11A).

Q11: **SELECT** Salary
 FROM EMPLOYEE;

Q11A: **SELECT** **DISTINCT** Salary
 FROM EMPLOYEE;



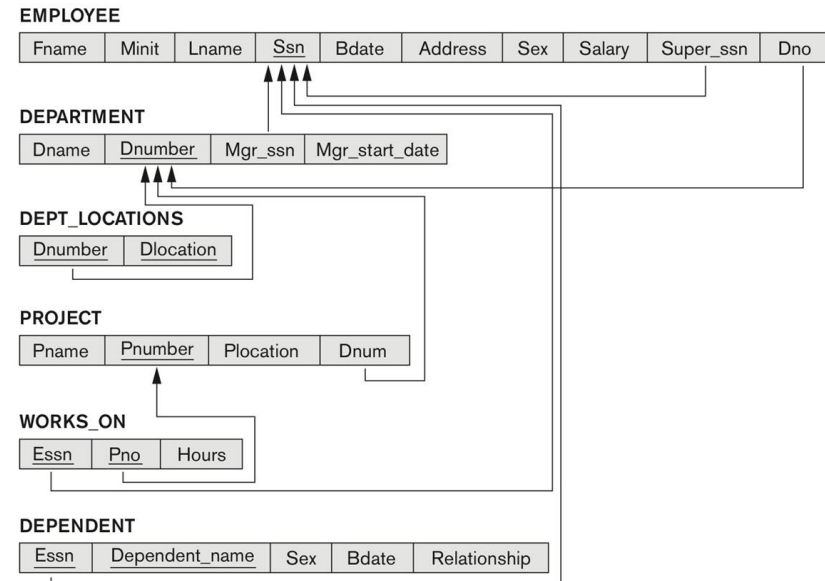
Set Operations

■ UNION, EXCEPT (difference), INTERSECT

Query 4. Make a list of all project numbers for projects that involve an employee whose last name is 'Smith', either as a worker or as a manager of the department that controls the project.

Q4A: (**SELECT** **DISTINCT** Pnumber
 FROM PROJECT, DEPARTMENT, EMPLOYEE
 WHERE Dnum=Dnumber **AND** Mgr_ssn=Ssn
 AND Lname='Smith')

 UNION
(**SELECT** **DISTINCT** Pnumber
 FROM PROJECT, WORKS_ON, EMPLOYEE
 WHERE Pnumber=Pno **AND** Essn=Ssn
 AND Lname='Smith');





String Matching (LIKE)

- **LIKE** is a comparison operator used for string pattern matching (similar to regular expressions).
 - **%** means an arbitrary number of zero or more characters
 - **_** (underscore) means a single character
- Examples:
 - **WHERE Address LIKE '%Houston, TX%';**
 - **WHERE Ssn LIKE '__ 1__ 8901';**



Arithmetic Operations

- Standard arithmetic operations (+, -, *, /) may be included as part of the **SELECT** clause.
- Show the resulting salaries if every employee working on the 'ProductX' project is given a 10 percent raise.

```
SELECT E.Fname, E.Lname, 1.1 * E.Salary AS Increased_sal  
FROM   EMPLOYEE AS E, WORKS_ON AS W, PROJECT AS P  
WHERE  E.Ssn=W.Essn AND W.Pno=P.Pnumber AND  
       P.Pname='ProductX';
```



Ordering of Results

- We can add the optional **ORDER BY** clause in order to sort query results in a certain order.
 - **ORDER BY** typically comes at the end of the query
 - **DESC**: descending order
 - **ASC**: ascending order

```
SELECT    <attribute list>
FROM      <table list>
[ WHERE   <condition> ]
[ ORDER BY <attribute list> ];
```

- **ORDER BY D.Dname DESC, E.Lname ASC, E.Fname ASC**



Examples

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Shero	shero@cs	18	3.2
53650	Shero	shero@math	19	3.8

Find all students
with age 18

```
SELECT *  
FROM Students S  
WHERE S.age=18
```

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Shero	shero@cs	18	3.2



Examples

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Shero	shero@cs	18	3.2
53650	Shero	shero@math	19	3.8

Names and logins of
all students
with age 18

```
SELECT S.name, S.login  
FROM Students S  
WHERE S.age=18
```

name	login
Jones	jones@cs
Shero	shero@cs



Examples

Enrolled

sid	cid	grade
53831	Carnatic101	C
53831	Reggae203	B
53650	Topology112	A
53666	History105	B

Names and cids of
students who received
A from one or more of
their courses

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Shero	shero@cs	18	3.2
53650	Shero	shero@math	19	3.8

```
SELECT S.name, E.cid  
FROM Students S, Enrolled E  
WHERE S.sid=E.sid AND E.grade="A"
```

S.name	E.cid
Shero	Topology112



Insertion

- **INSERT**: add tuples to a relation
 - **INSERT INTO** *<table_name>* **VALUES** (...)
 - Attribute order should be same as CREATE TABLE
 - Constraints wrt data types (INT, CHAR, ..) and integrity (related to PK, FK, ..) are enforced

U1: **INSERT INTO** **EMPLOYEE**
 VALUES ('Richard', 'K', 'Marini', '653298653', '1962-12-30', '98
 Oak Forest, Katy, TX', 'M', 37000, '653298653', 4);

CREATE TABLE EMPLOYEE			
(Fname	VARCHAR(15)		NOT NULL,
Minit	CHAR,		
Lname	VARCHAR(15)		NOT NULL,
Ssn	CHAR(9)		NOT NULL,
Bdate	DATE,		
Address	VARCHAR(30),		
Sex	CHAR,		
Salary	DECIMAL(10,2),		
Super_ssn	CHAR(9),		
Dno	INT		NOT NULL,
PRIMARY KEY (Ssn),			



Deletion

- **DELETE**: remove tuples from a relation
 - **DELETE FROM** *<table_name>* **WHERE** (...)
 - The WHERE clause indicates the condition regarding which tuples will be deleted

U4A:	DELETE FROM	EMPLOYEE
	WHERE	Lname='Brown';
U4B:	DELETE FROM	EMPLOYEE
	WHERE	Ssn='123456789';
U4C:	DELETE FROM	EMPLOYEE
	WHERE	Dno=5;
U4D:	DELETE FROM	EMPLOYEE;



Update

- **UPDATE**: modify attribute values of tuples
 - **UPDATE** *<table_name>* **SET** *<changes>* **WHERE** (...)

Change the location and controlling department number of project #10 to "Bellaire" and 5, respectively.

```
UPDATE      PROJECT
SET          PLOCATION = 'Bellaire',
              DNUM = 5
WHERE        PNUMBER=10
```

Give all employees in the "Research" department a 10% raise in salary.

```
UPDATE      EMPLOYEE
SET          SALARY = SALARY *1.1
WHERE        DNO IN  ( SELECT DNUMBER
                       FROM DEPARTMENT
                       WHERE DNAME='Research' )
```



Examples

INSERT INTO
VALUES

Students
(53688, 'Smith', 'smith@ee', 18, 3.2)

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@ee	18	3.2

DELETE FROM Students S
WHERE S.name = 'Jones'

DROP TABLE Students



Destroy the Students relation.
Schema and tuples are deleted.



Exercises

- Consider the following schema:
 - Sailors (sid, sname, rating, age)
 - Boats (bid, bname, color)
 - Reserves (sid, bid, day)

<i>Sailors</i>	<u>sid</u>	sname	rating	age
	22	dustin	7	45.0
	31	lubber	8	55.5
	58	rusty	10	35.0

<i>Boats</i>	<u>bid</u>	bname	color
	101	Interlake	Blue
	102	Interlake	Red
	103	Clipper	Green
	104	Marine	Red

<i>Reserves</i>	<u>sid</u>	<u>bid</u>	<u>day</u>
	22	101	10/10/96
	58	103	11/12/96



Exercises

- Sailors (sid, sname, rating, age)
- Boats (bid, bname, color)
- Reserves (sid, bid, day)
- Find the names of sailors who reserved boat number 103.

```
SELECT sname
FROM   Sailors, Reserves
WHERE  Sailors.sid=Reserves.sid AND bid=103
```

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid=R.sid AND R.bid=103
```



Exercises

- Sailors (sid, sname, rating, age)
- Boats (bid, bname, color)
- Reserves (sid, bid, day)
- For sailors whose names begin and end with letter B and contain at least three characters, display their name and twice their age.

```
SELECT  S.sname, 2*S.age
FROM    Sailors S
WHERE   S.sname LIKE 'B_%B'
```



Exercises

- Sailors (sid, sname, rating, age)
- Boats (bid, bname, color)
- Reserves (sid, bid, day)
- Find the IDs of sailors who have reserved a red boat or a green boat.

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid
      AND (B.color='red' OR B.color='green')
```

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid
      AND B.color='red'
UNION
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid
      AND B.color='green'
```




Exercises

- Sailors (sid, sname, rating, age)
- Boats (bid, bname, color)
- Reserves (sid, bid, day)
- Find the IDs of sailors who have reserved a red boat but never reserved a green boat.

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid
      AND B.color='red'
EXCEPT
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid
      AND B.color='green'
```



Exercises

- Sailors (sid, sname, rating, age)
- Boats (bid, bname, color)
- Reserves (sid, bid, day)
- Find the IDs of sailors who have reserved a red boat and a green boat.

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid
      AND B.color='red'
INTERSECT
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid
      AND B.color='green'
```

```
SELECT R1.sid
FROM Boats B1, Reserves R1,
      Boats B2, Reserves R2
WHERE R1.sid = R2.sid AND
      R1.bid=B1.bid AND R2.bid=B2.bid
      AND (B1.color='red' AND B2.color='green')
```



Conclusion

- SQL is a comprehensive language for relational database management (select, insert, modify, delete, ...)
- So far, we have covered fundamental SQL commands
 - Many advanced aspects and functionalities exist (we will cover some more, but not all)
 - You'll learn about some of them in the PSs + HWs
- **Next topic: Advanced SQL**
 - Nested SQL queries
 - IN, EXISTS, NOT EXISTS, ...
 - GROUP BY - HAVING
 - Aggregates: COUNT, SUM, MIN, MAX, ...
 - Views