



**KOÇ
UNIVERSITY**

Database Management Systems

Transaction Management

M. Emre Gürsoy

Assistant Professor
Department of Computer Engineering

www.memregursoy.com




Motivation

- Consider the database of a **banking application**
 - Multiple concurrent users and processes: withdrawals from ATM, mobile payments, money transfers...
- Consider that **Peter** wants to send 100 TL to **Stewie**
 - Read money in Peter's account (**a**)
 - **$a = a - 100$** ; write it to the database
 - Read money in Stewie's account (**b**)
 - **$b = b + 100$** ; write it to the database
- Simultaneously, the **Griffin Family Fund** wants to calculate the total money **Peter Griffin** and **Stewie Griffin** have



400 TL




100 TL



300 TL

100 TL




100 TL



300 TL

100 TL



200 TL

Griffin
Family
Fund


500 TL

Things are fine if:

- Money is transferred
- Afterwards, the sum is calculated



400 TL




100 TL



300 TL

100 TL



100 TL

Griffin
Family
Fund


400 TL

**But what if the concurrent
transaction is scheduled right
before the money is credited
to Stewie's account?**



300 TL

100 TL



100 TL




SYSTEM CRASH

How much money should be
in their account after the
system is restored?



Transactions

- A **transaction** is the execution of a sequence of one or more operations (e.g., SQL queries) on a DB to perform some **higher-level function**.
 - Money transfer is a higher-level function.
 - Summation of funds is a higher-level function.
- Transaction = **one logical unit of work**
- Technically, a transaction contains a sequence of **reads** and **writes**: R(A), W(A), R(B), W(C), ...
 - A,B,C: database objects
 - DBMS's abstract view of a high-level program/function
- A transaction starts with **BEGIN**
- A transaction ends with **COMMIT** or **ABORT** (aka **ROLLBACK**)





Strawman Approach

- Say that we have two transactions:
 - **T1**: Peter -> Stewie money transfer
 - **T2**: Griffin Family Fund summation
 - Execute T1 first, then T2
- **The strawman approach:**
 - In general, when you have N transactions, execute them in **serial order** (one-by-one).
 - Before each transaction BEGINS, copy the whole DB to a new file and make all changes on that file.
 - If transaction finishes successfully and COMMITs, copy the file's contents to the DB.
 - If transactions ABORT, just delete the dirty copy.



Strawman Approach

- What's bad about the **strawman approach**?
 - No parallelism, no multi-threading
 - Some transactions **can** be parallelized
 - **T1**: R(A), W(A) **T2**: R(B), W(B)
 - Low throughput and increased response times to users
 - Copying large volumes of data back and forth
- We certainly want to do things smarter than the strawman approach, but...
 - We want our results to be correct
 - We want our data to be safe




ACID Properties

- **Atomicity**
 - "All or nothing": Either all actions in a transaction happen, or none happen.
- **Consistency**
 - If the DB is initially consistent, after a valid transaction, the DB should again be consistent.
- **Isolation**
 - "As if alone": Execution of one transaction is isolated from the executions of other transactions.
- **Durability**
 - "The results stay": If a transaction commits, its effects persist (despite OS crashes, electricity outages, ...).



Atomicity

- From the user's point of view: transactions always execute either all actions, or execute no actions at all.



- Approach #1: **Logging**
 - Log all actions, undo the actions of aborted txns
 - Used in almost every DBMS
 - Also good for auditing
- Approach #2: **Shadow Paging**
 - DBMS makes copies of **pages** and txns make changes to those copies. Only when the txn commits, the page is made visible to others.
 - Very few DBMS systems do this (it's quite slow)



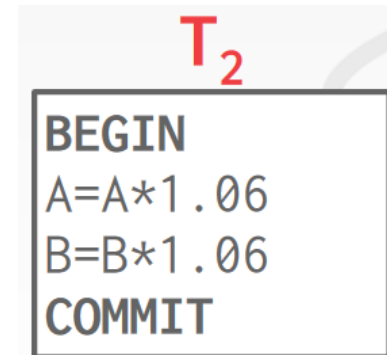
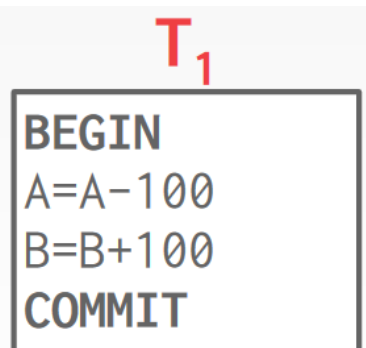
Isolation

- Users submit transactions, and each transaction executes as if it was running by itself.
- We can achieve this using the **strawman approach**: run transactions one-by-one, in serial order.
 - But that's inefficient
- The DBMS achieves **concurrency** by **interleaving** the actions (reads/writes of DB objects) of transactions.
 - We need a way to interleave transactions but still make it appear as if they ran one at a time.
- The task of finding a proper interleaving of operations from multiple transactions is achieved by **scheduling** and **concurrency control**.



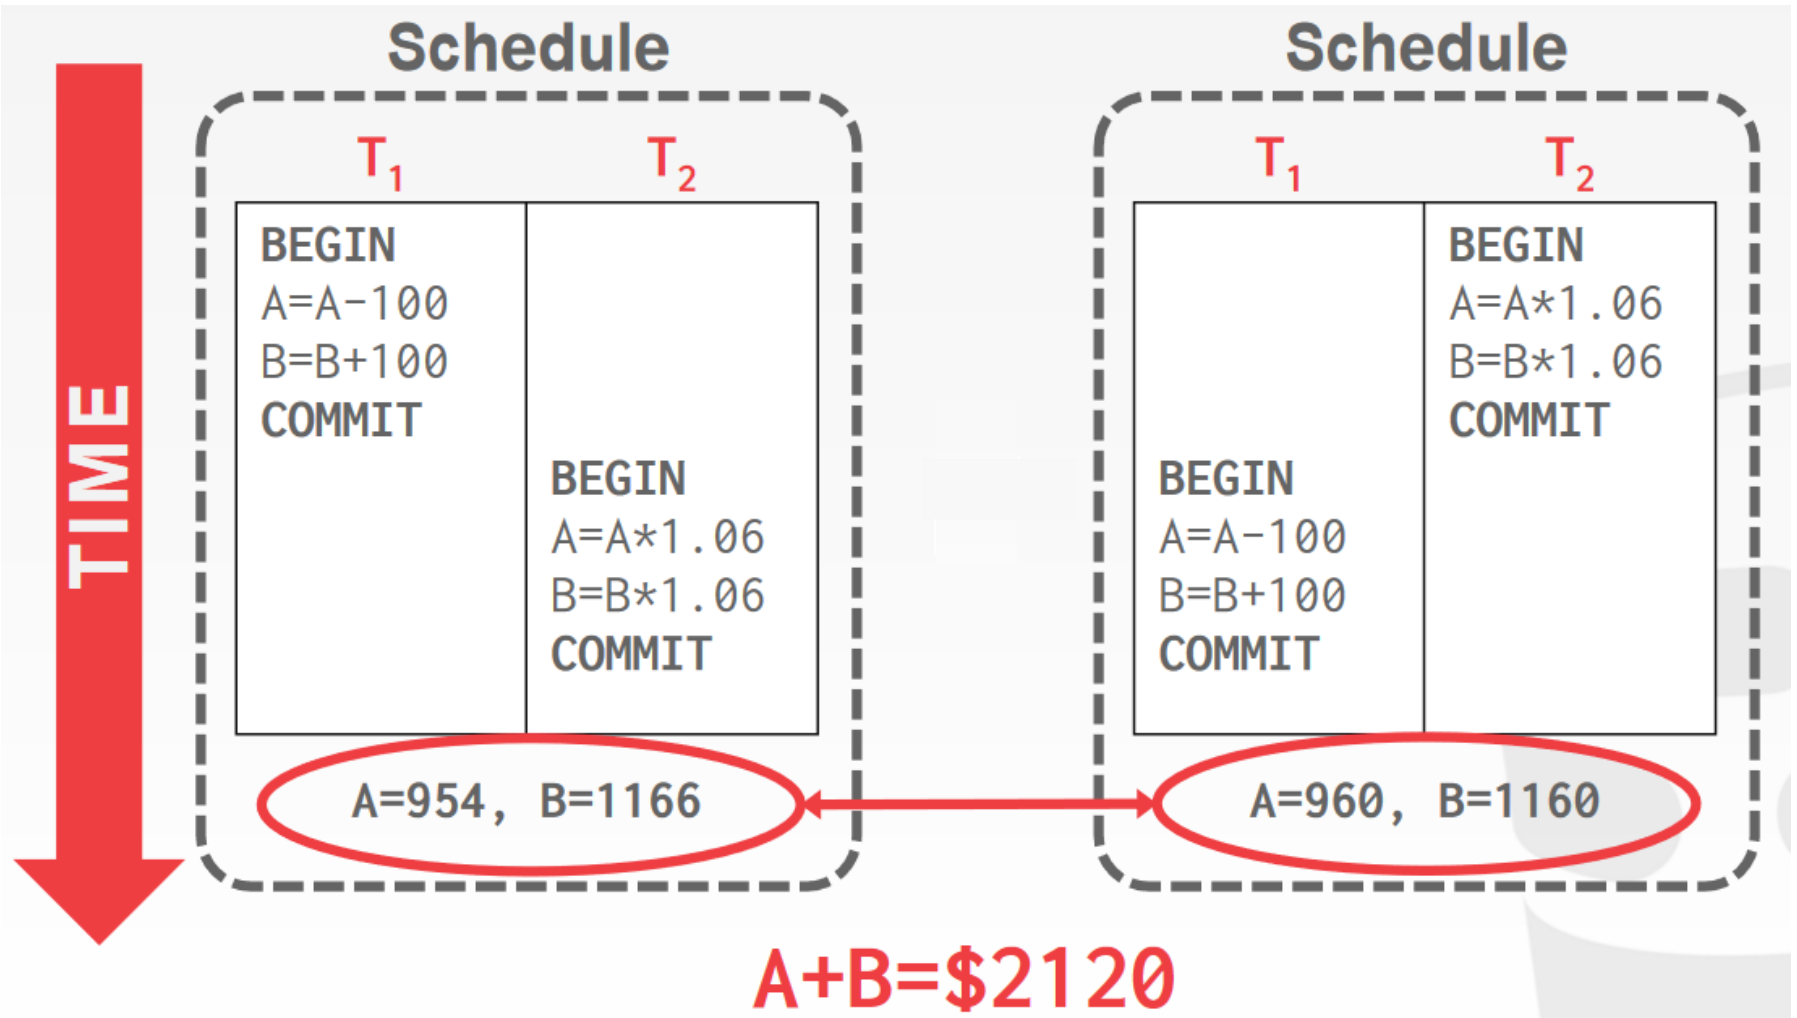
Scheduling Example

- Assume at first **A** and **B** each have \$1000.
- Say that **T₁** and **T₂** are submitted together (there's no guarantee which one executes first).
- What are the legal outcomes?
 - A** = 954, **B** = 1166 → **A+B** = 2120
 - A** = 960, **B** = 1160 → **A+B** = 2120
- Regardless of which executes first, **A+B** should be 2120.
 - "The net effect is equal to **some** serial execution of T1 and T2".



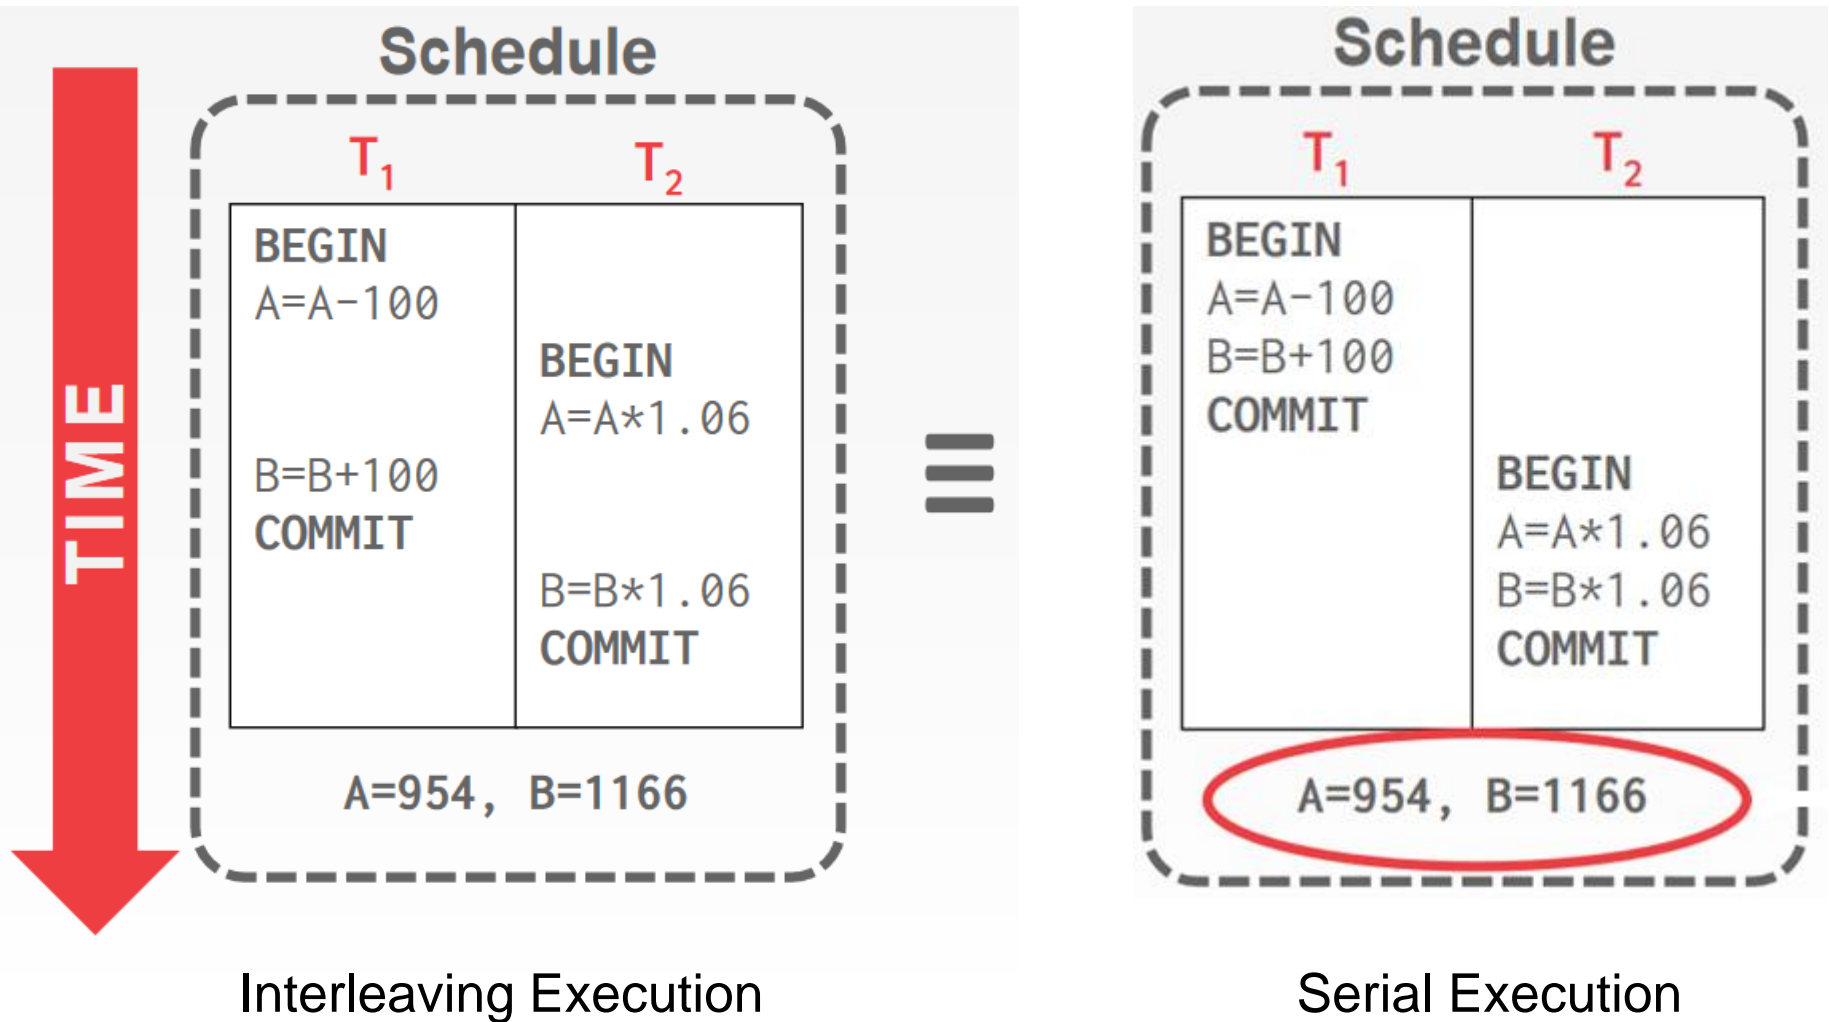


Serial Execution



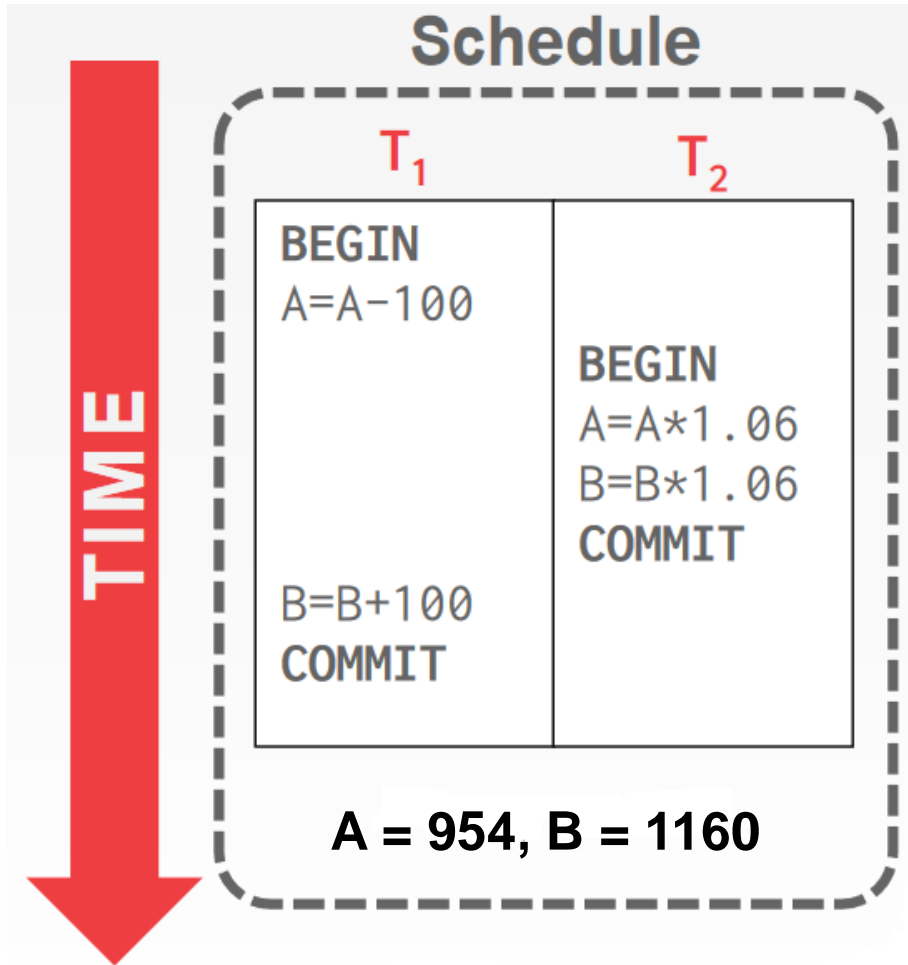


Interleaving Execution (Good)





Interleaving Execution (Bad)

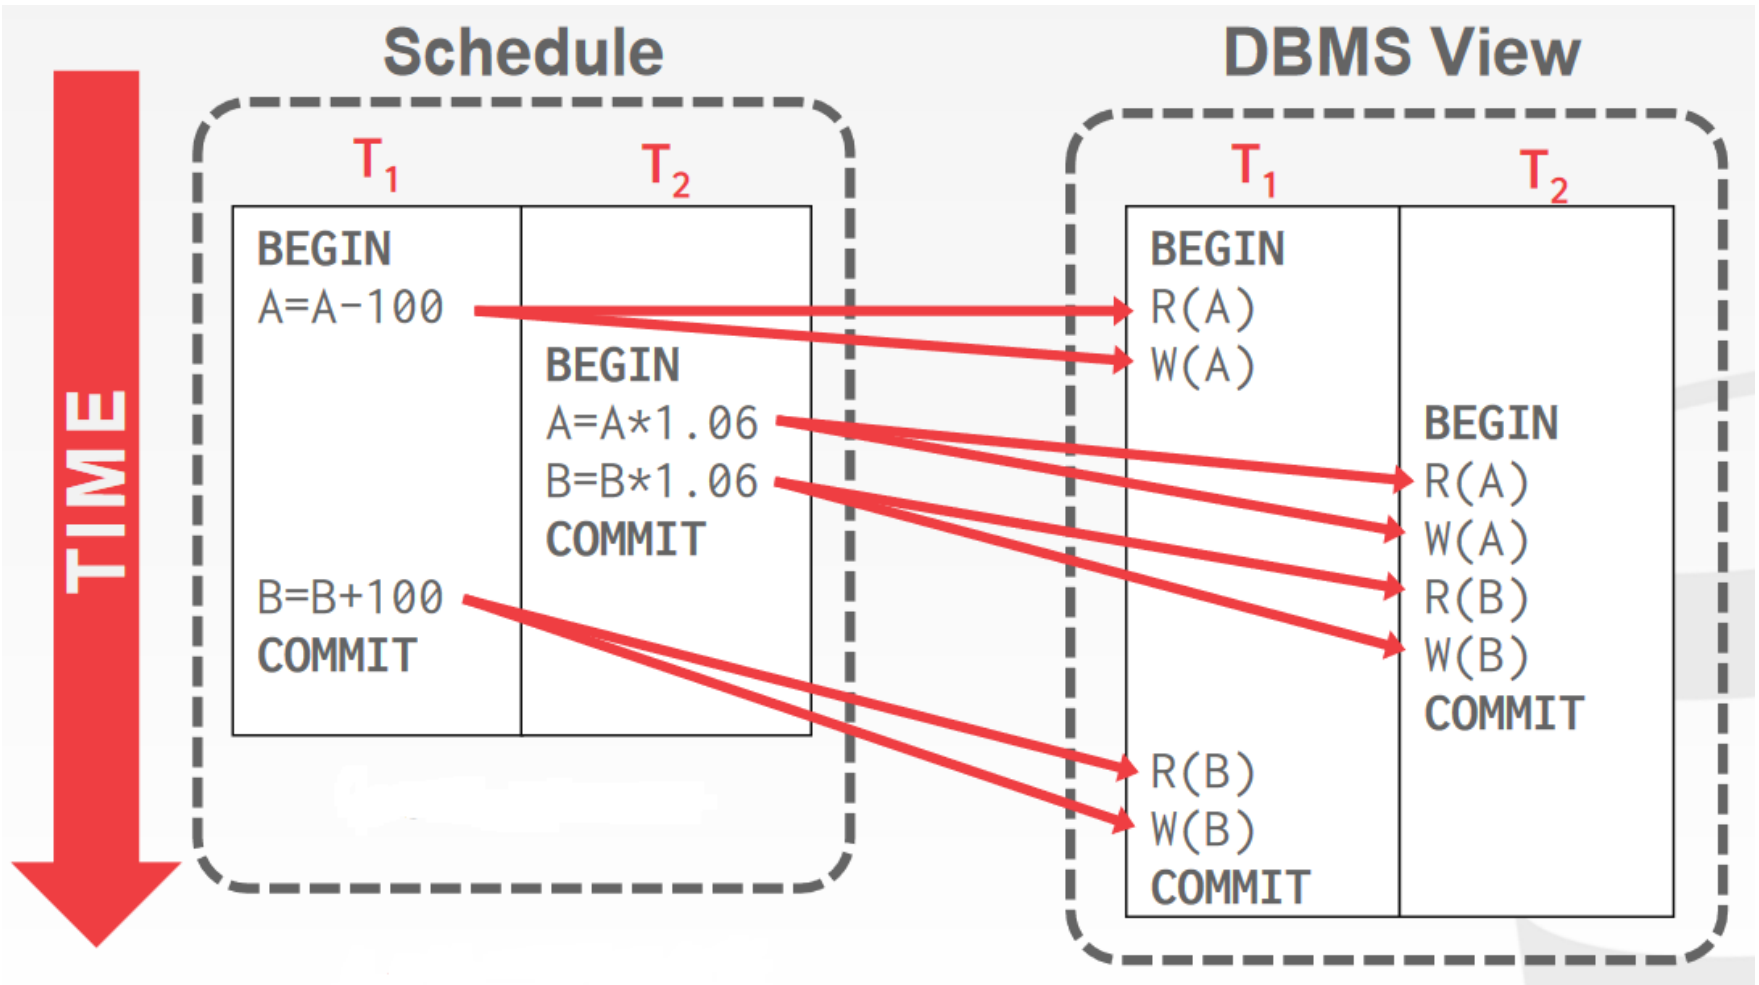


\neq $A=954, B=1166$
or
 $A=960, B=1160$

$A+B = \$2114$, the bank is missing \$6!



DBMS Perspective





Schedules

- **Serial schedule:** A schedule that does not interleave the actions of different transactions.
- **Equivalent schedules:** Two schedules are **equivalent** if the effect of executing the first schedule is identical to the effect of executing the second schedule.
- **Serializable schedule:** A schedule that is equivalent to **some** serial execution of the transactions.
 - If each transaction preserves consistency, every serializable schedule preserves consistency.



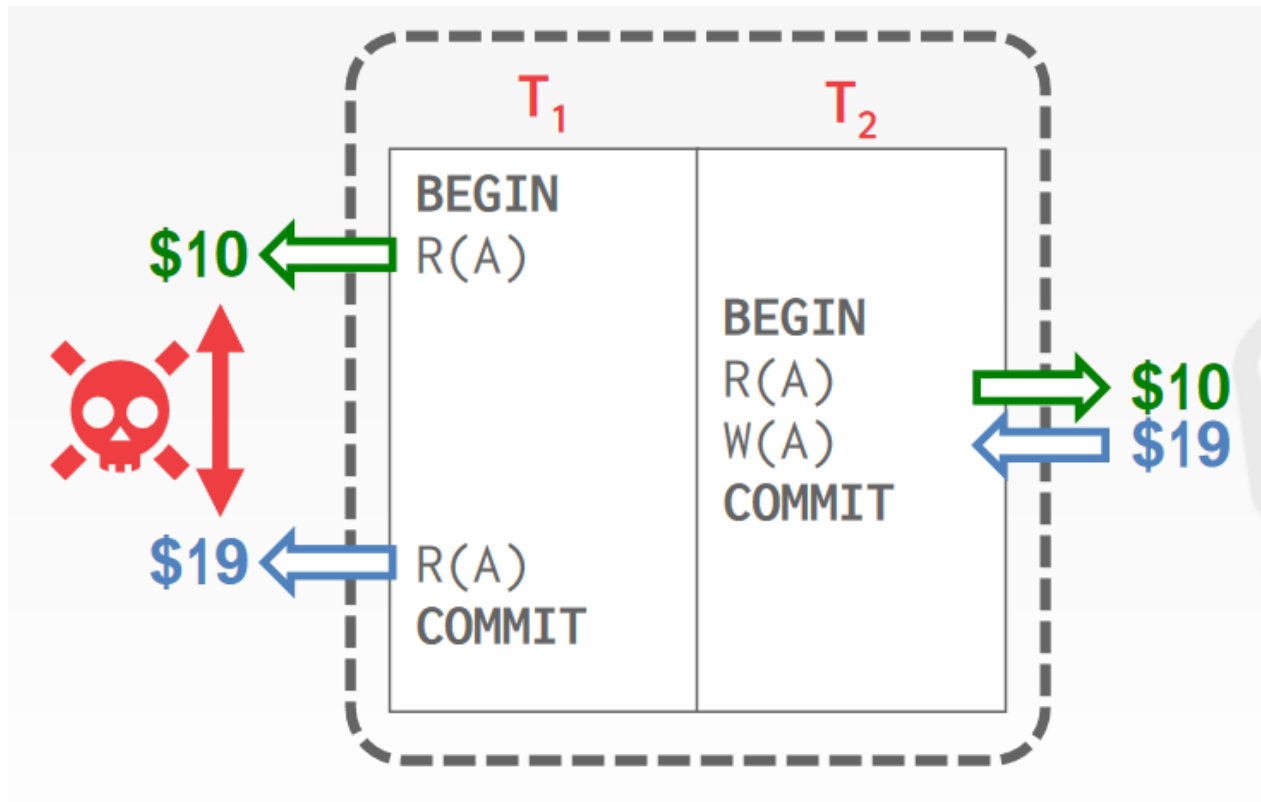
Conflicting Operations

- Two operations **conflict** if:
 - They are by **different transactions**
 - They are on the **same object** and at least one of them is a **write** operation
 - **Why do multiple reads not cause a problem?**
- Types of **problems** that conflicts can lead to:
 - Read-Write Problems **(R-W)**
 - Write-Read Problems **(W-R)**
 - Write-Write Problems **(W-W)**



R-W Problems

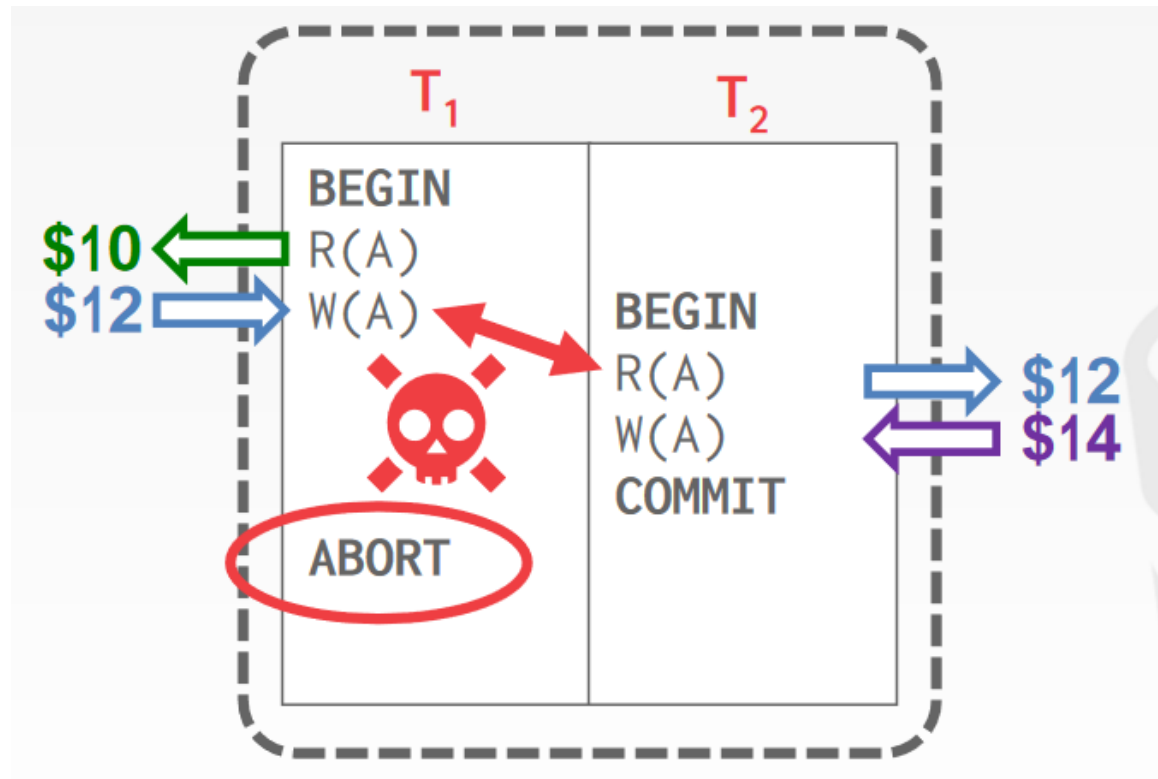
- "Unrepeatable reads"





W-R Problems

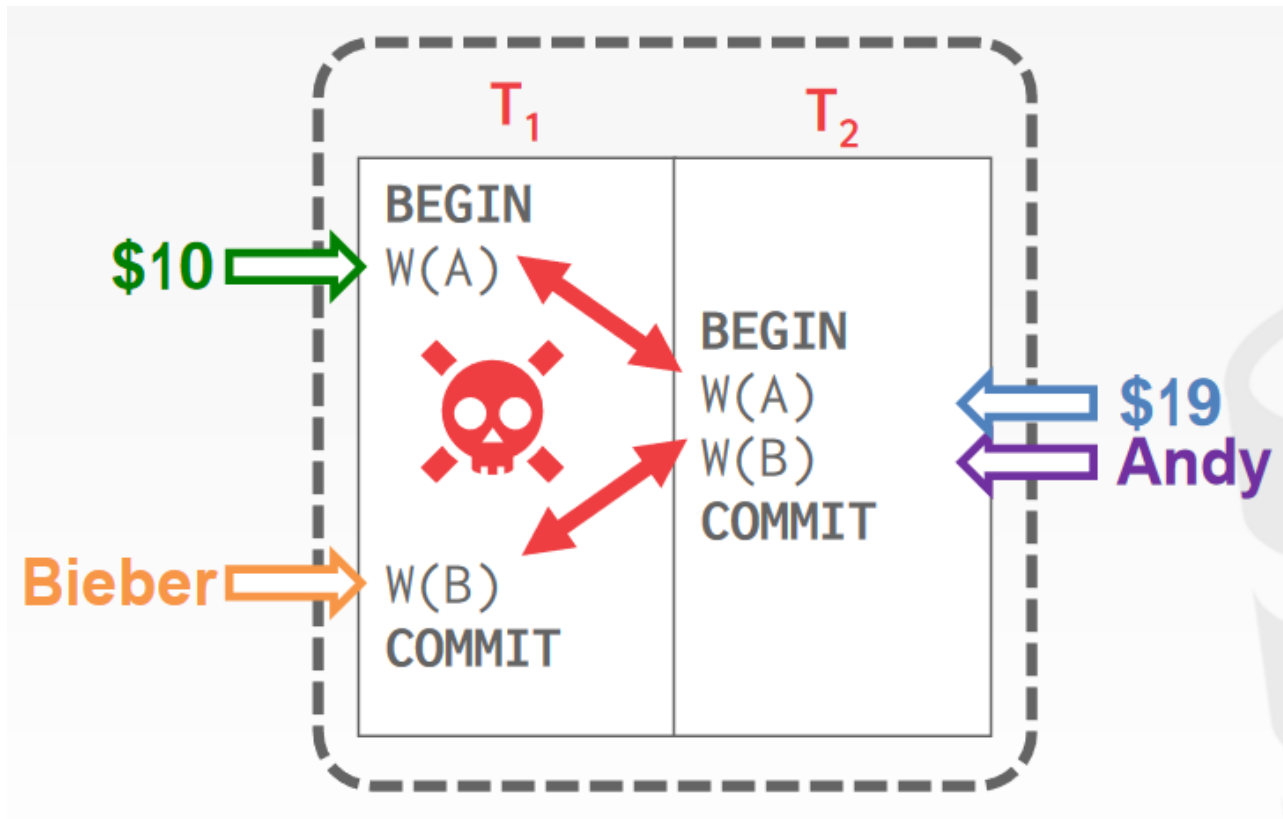
- "Dirty reads": reading uncommitted data





W-W Problems

- Overwriting uncommitted data



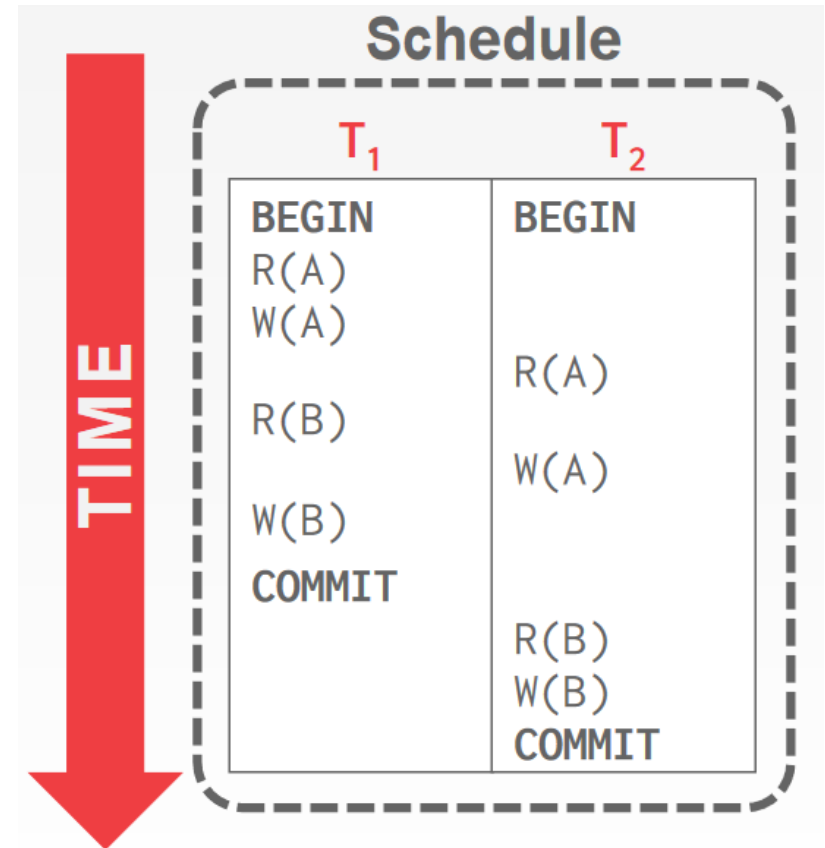
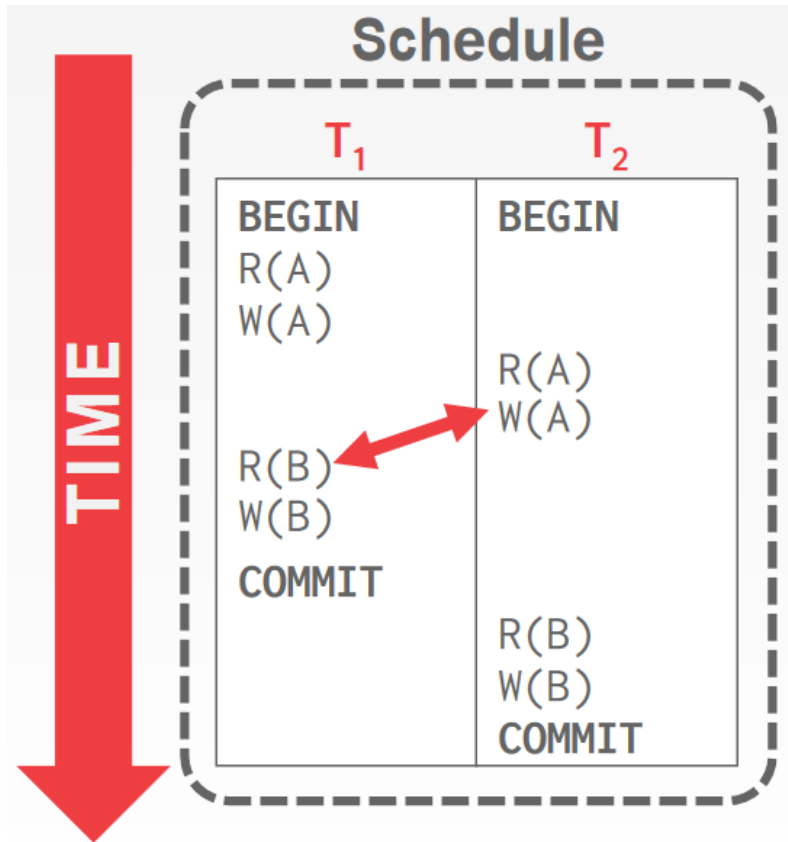


Conflict Serializability

- Two schedules are **conflict equivalent** iff:
 - They involve the same actions of the same transactions
 - Every pair of conflicting actions is ordered the same way
- Schedule **S** is **conflict serializable** if:
 - S** is conflict equivalent to some serial schedule
- Swapping method**: to check conflict serializability
 - Schedule **S** is conflict serializable if you can transform **S** into a serial schedule by swapping the order of consecutive non-conflicting operations of different transactions.

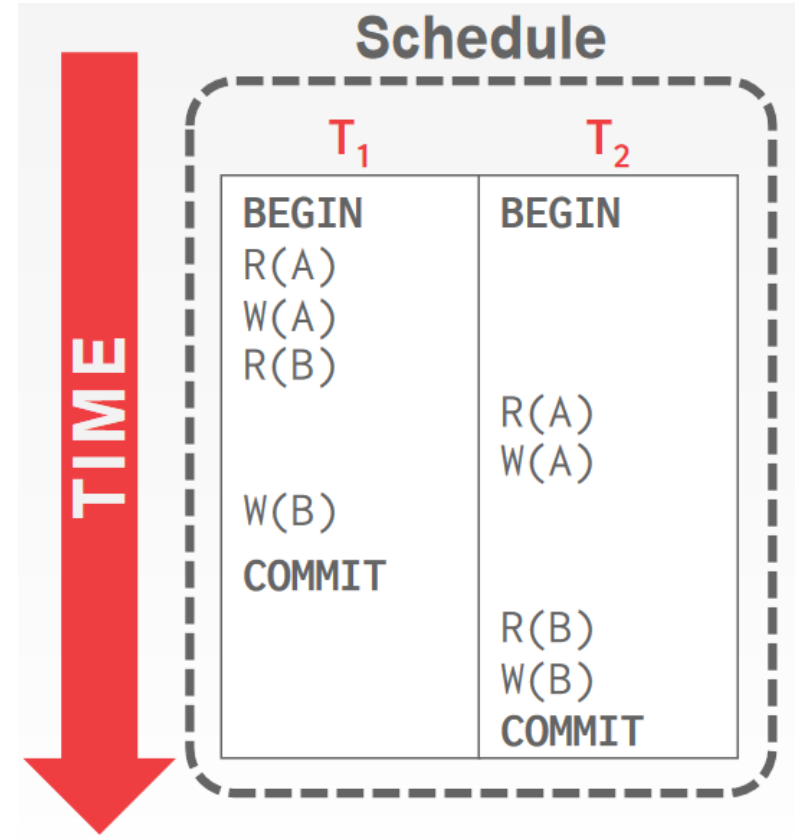
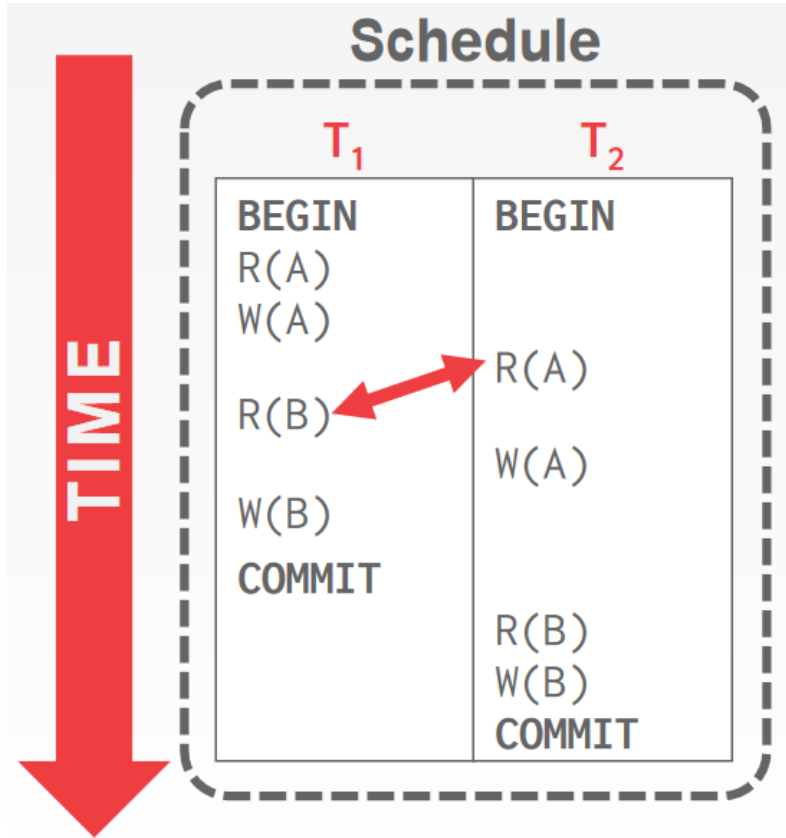


Swapping Method



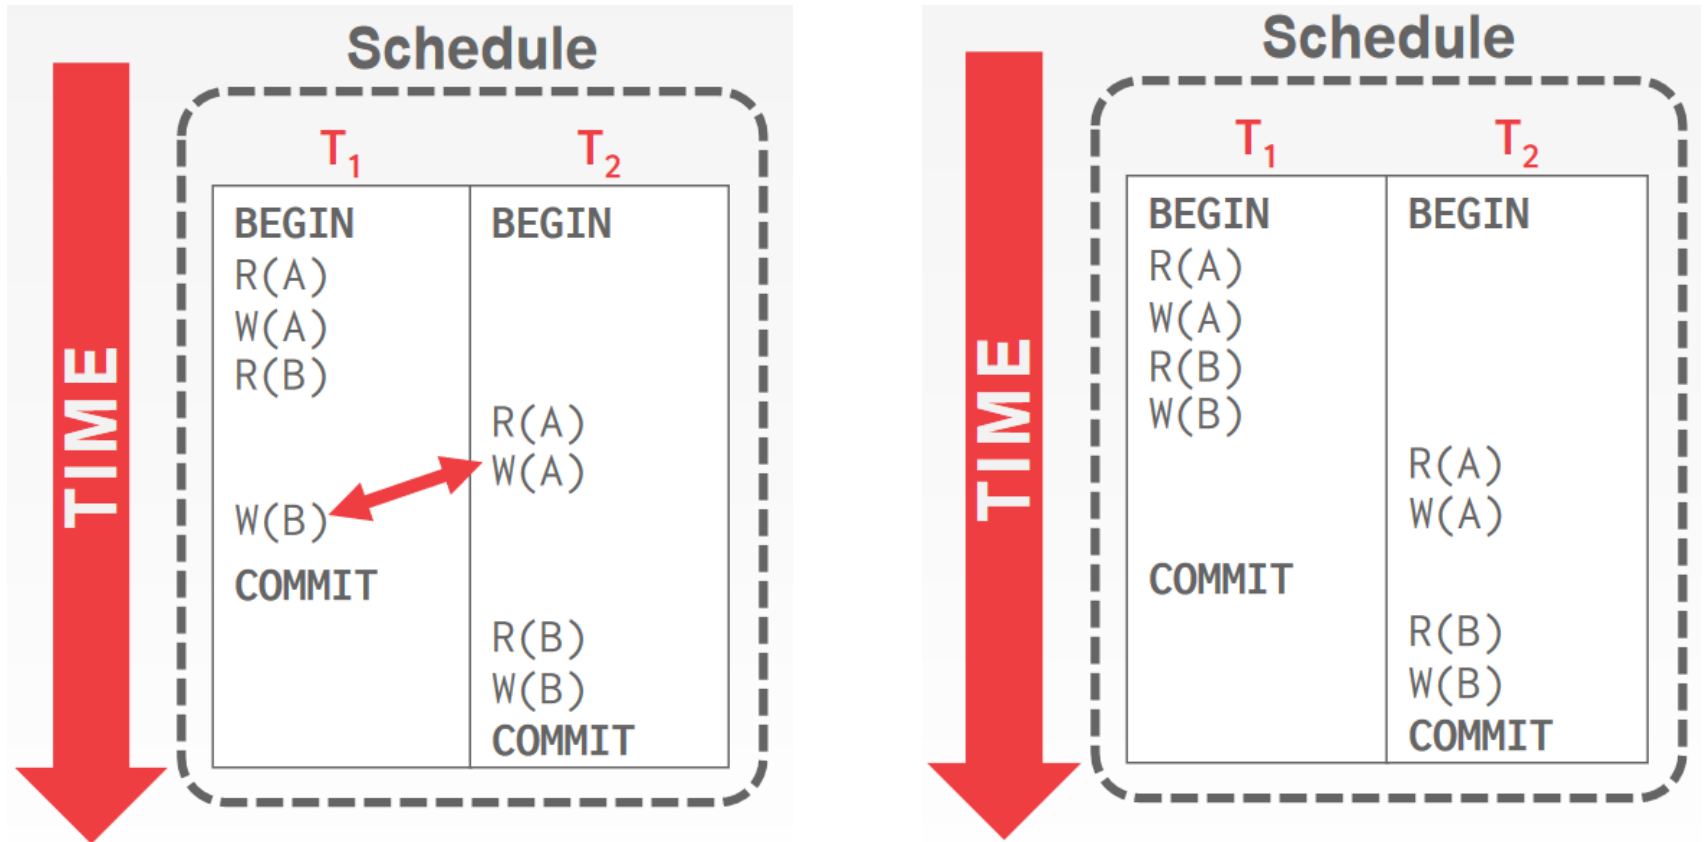


Swapping Method





Swapping Method

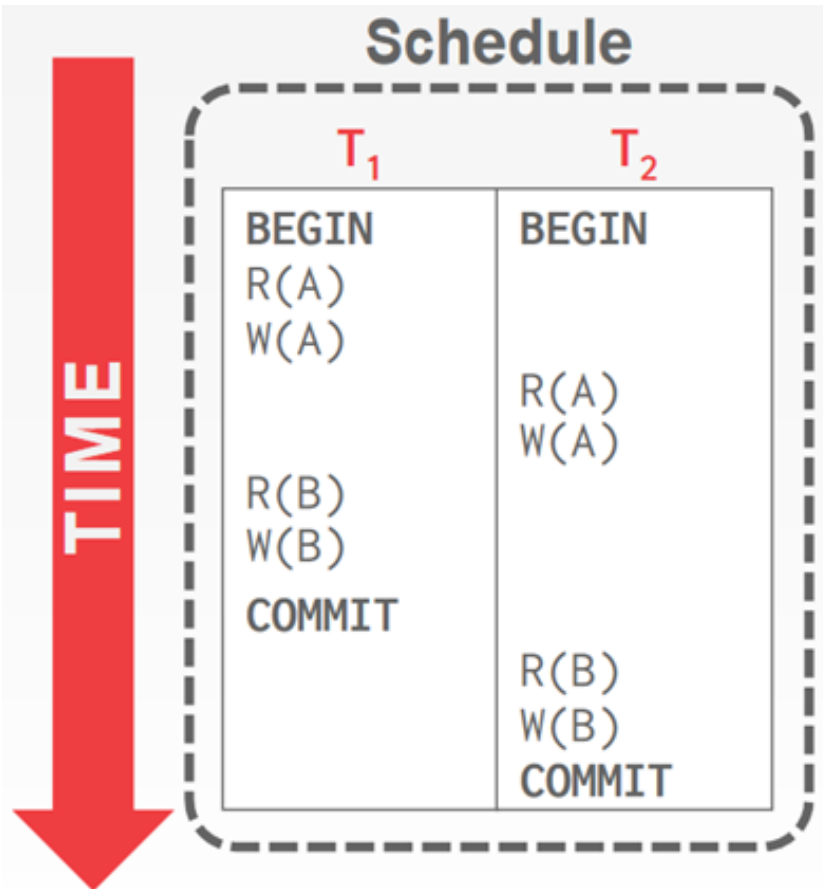


Fast-forward one more step: comparing W(B) and R(A)

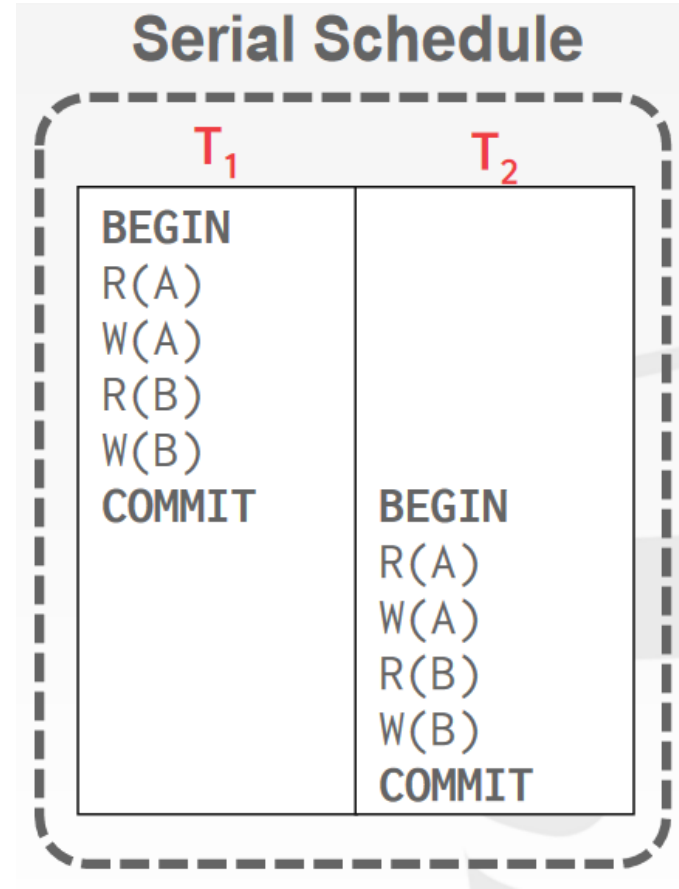


Swapping Method

Original Schedule



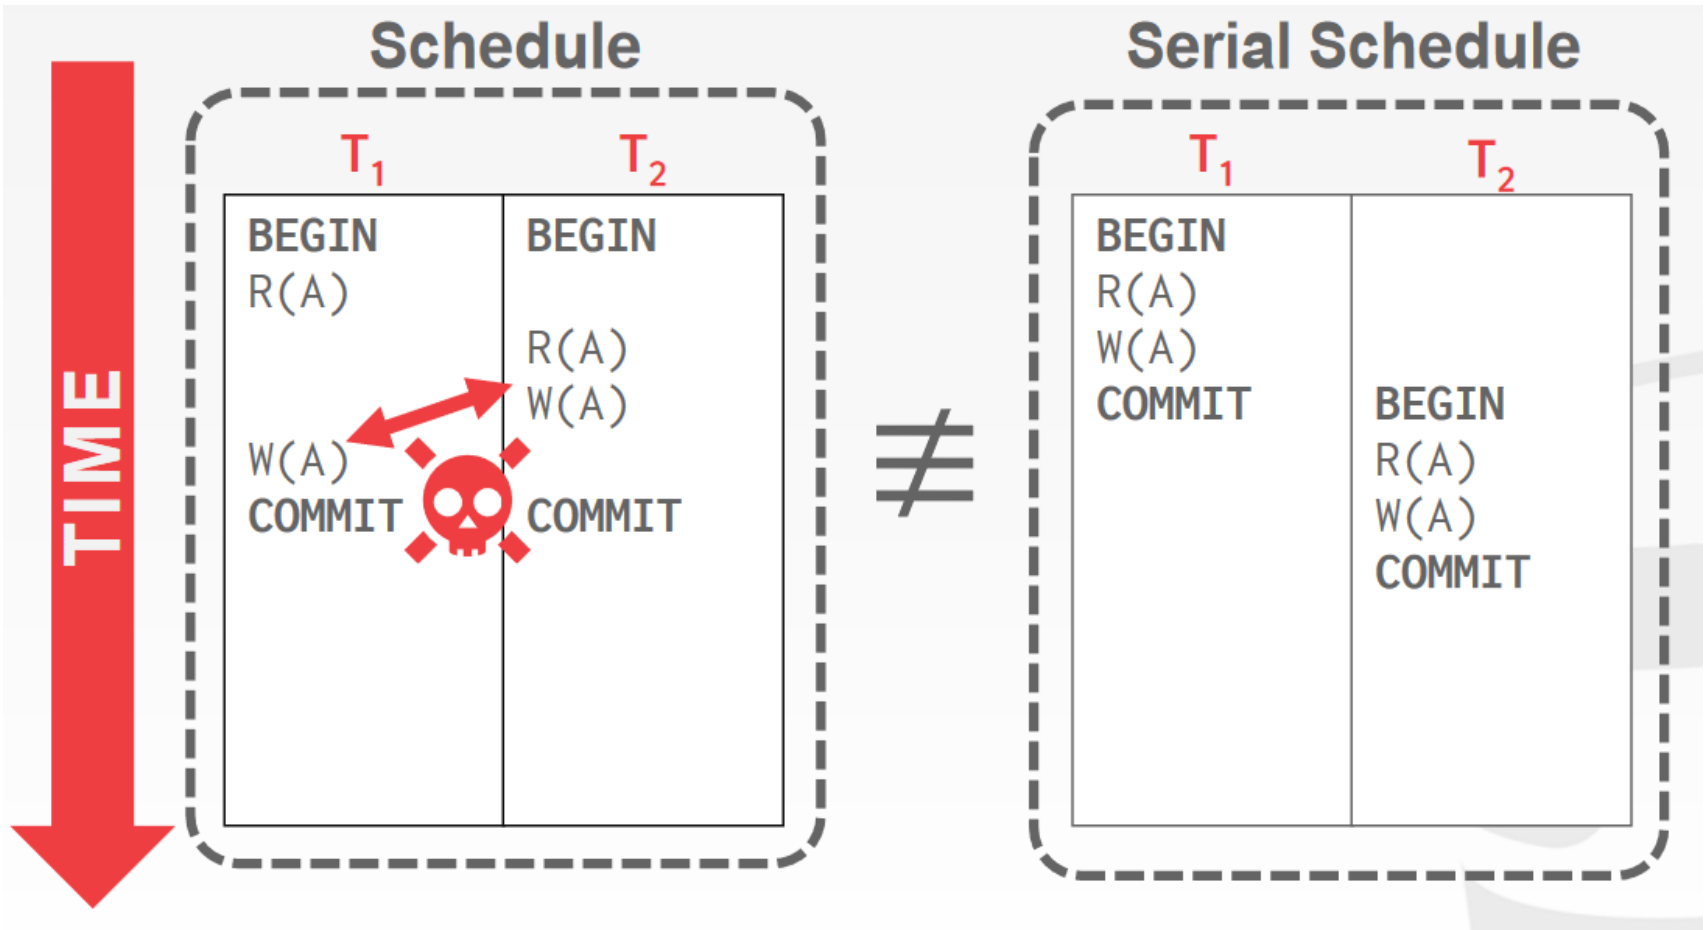
Transformed into:



Conflict serializable!



Swapping Method






Conflict Serializability

- Swapping method (i.e., checking conflict serializability by moving actions up and down) is feasible when:
 - There are **few transactions** in the schedule
 - Each transaction contains **few actions**
- But very cumbersome in **complex scenarios!**
 - Tens of transactions, hundreds of objects, thousands of reads & writes
- How to solve this problem better?
 - **DEPENDENCY GRAPH**



Dependency Graphs


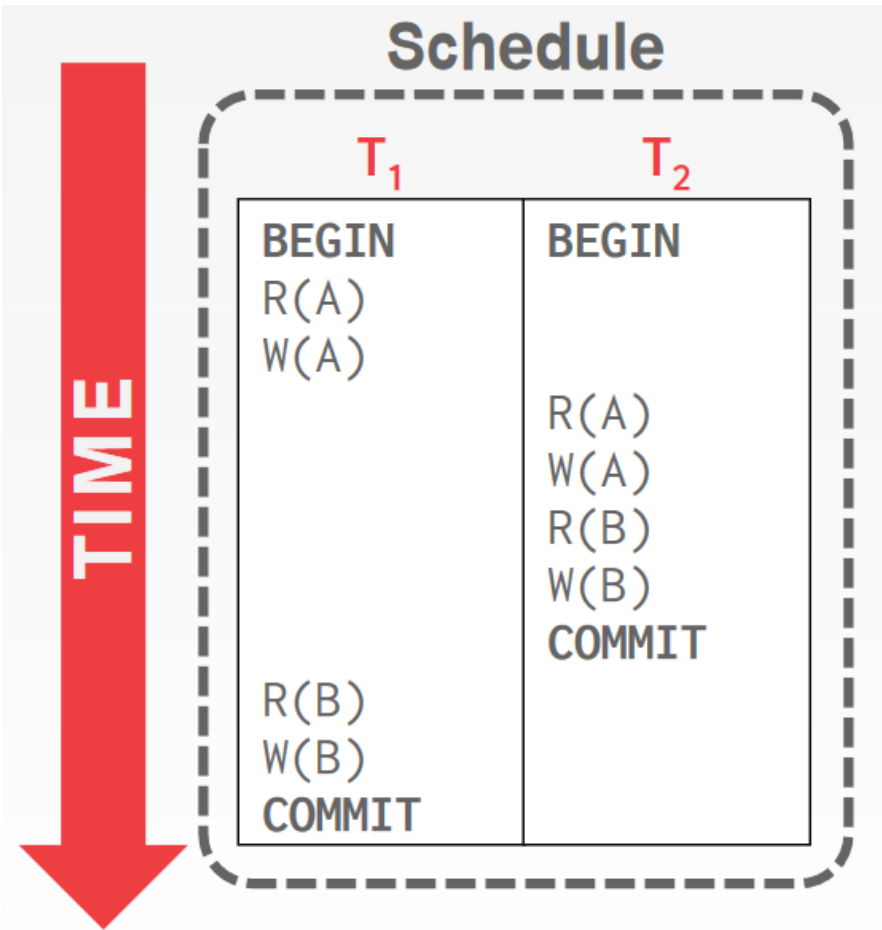
- One node per transaction
- Edge from T_i to T_j if:
 - An action A_i of T_i conflicts with an action A_j of T_j
 - A_i appears earlier than A_j in the schedule



- Theorem: **A schedule is conflict serializable if and only if its dependency graph is acyclic.**


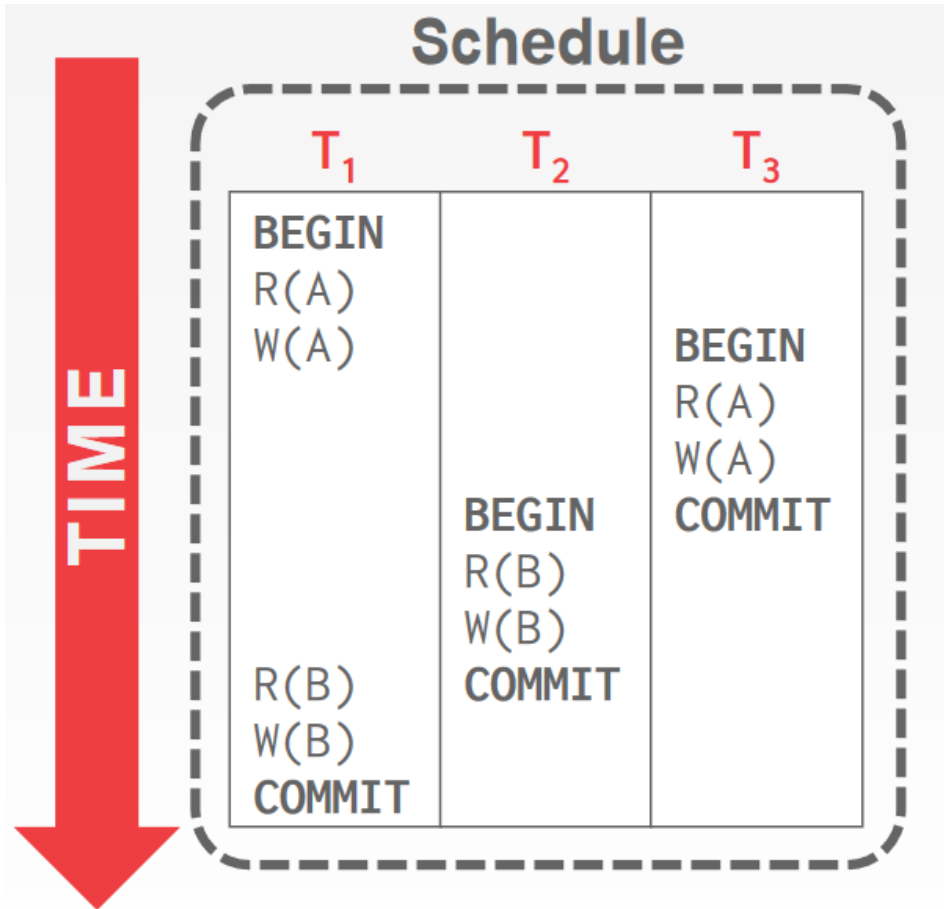


Example #1





Example #2



Which serial execution is this equivalent to?

T_2, T_1, T_3



Scheduling – Informal Summary

- **Serial** schedule: transactions executed one-by-one
- Two schedules are **equivalent** if their net effect is the same
- A schedule (consisting of multiple transactions) is **serializable** if it is equivalent to some serial execution of the transactions
- Schedules may contain **conflicts** (conflicting actions)
- Two schedules are **conflict equivalent** if their conflicting actions are ordered the same way
- A schedule is **conflict serializable** if it is conflict equivalent to some serial execution of the transactions
- We check conflict serializability using the **swapping method** and/or the **dependency graph**